

Web 2.0

Lecture 4: HATEOAS, Scalability and Description

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <http://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <http://vitvar.com/courses/w20>



EVROPSKÁ
UNIE

Evropský sociální fond
Praha & EU: Investujeme do vaší budoucnosti

Modified: Mon Mar 20 2017, 20:22:04
Humla v0.3

REST Core Principles

- REST architectural style defines constraints
 - *if you follow them, they help you to achieve a good design, interoperability and scalability.*
- Constraints
 - *Client/Server*
 - *Statelessness*
 - *Cacheability*
 - *Layered system*
 - *Uniform interface*
- Guiding principles
 - *Identification of resources*
 - *Representations of resources and self-descriptive messages*
 - *Hypermedia as the engine of application state (HATEOAS)*

Overview

- HATEOAS
 - *Stateful vs. Stateless*
 - *Links and Preconditions*
- Scalability

HATEOAS

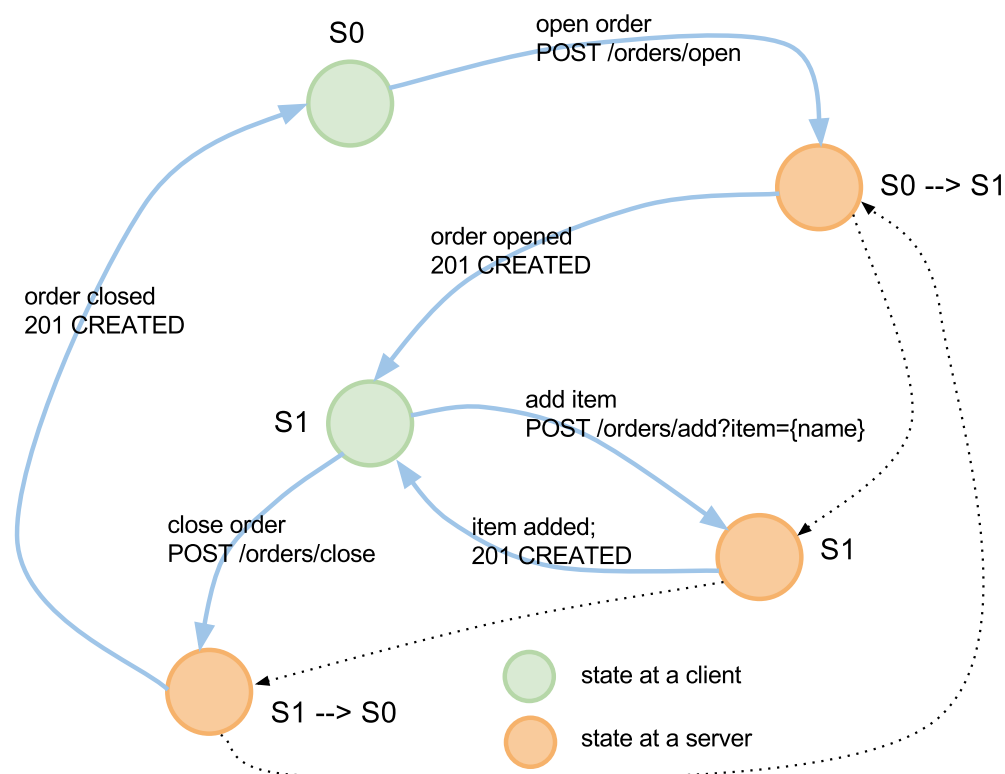
- HATEOAS = Hypertext as the Engine for Application State
 - *The REST core principle*
 - ***Hypertext***
 - *Hypertext is a representation of a resource with **links***
 - *A link is an URI of a resource*
 - *Applying an access to a resource via its link = state transition*
- Statelessness
 - *A service does not use a memory to remember a state*
 - *HATEOAS enables stateless implementation of services*

Overview

- HATEOAS
 - *Stateful vs. Stateless*
 - *Links and Preconditions*
- Scalability

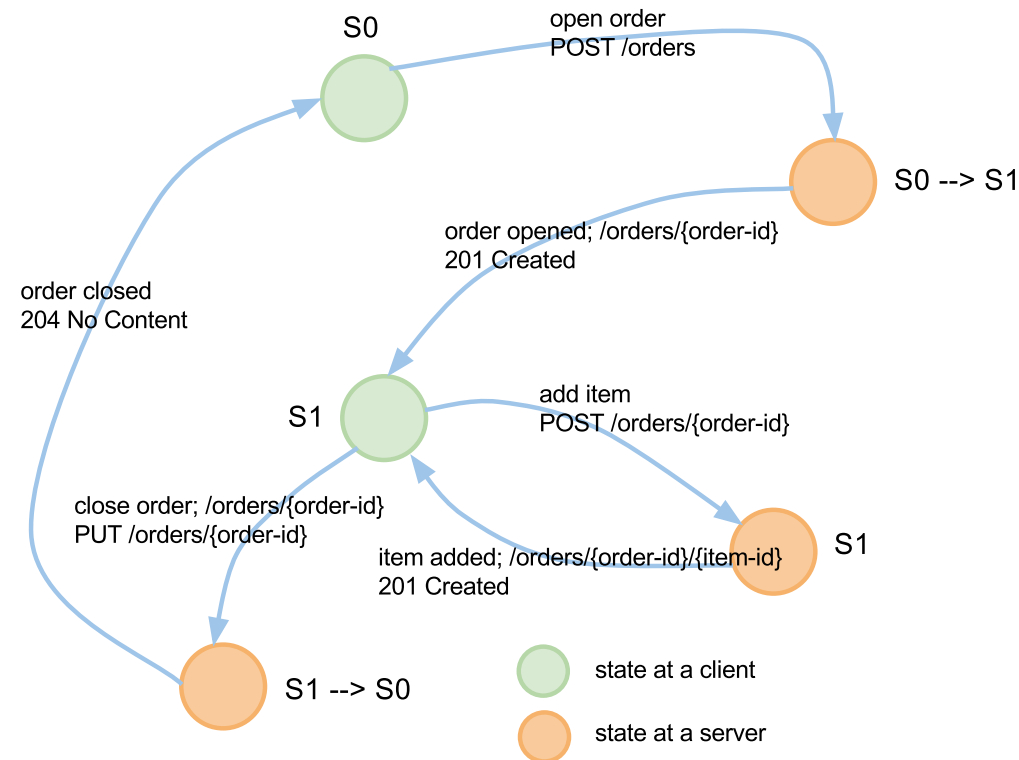
Stateful server

- Sessions to store the application state
 - Recall HTTP state management in MDW
 - The app uses a server memory to remember the state
 - when server restarts, the app state is lost



Stateless server

- HTTP and hypermedia to transfer the app state
 - *Does not use a server memory to remember the app state*
 - *State transferred between a client and a service via HTTP metadata and resources' representations*



Persistent Storage and Session Memory

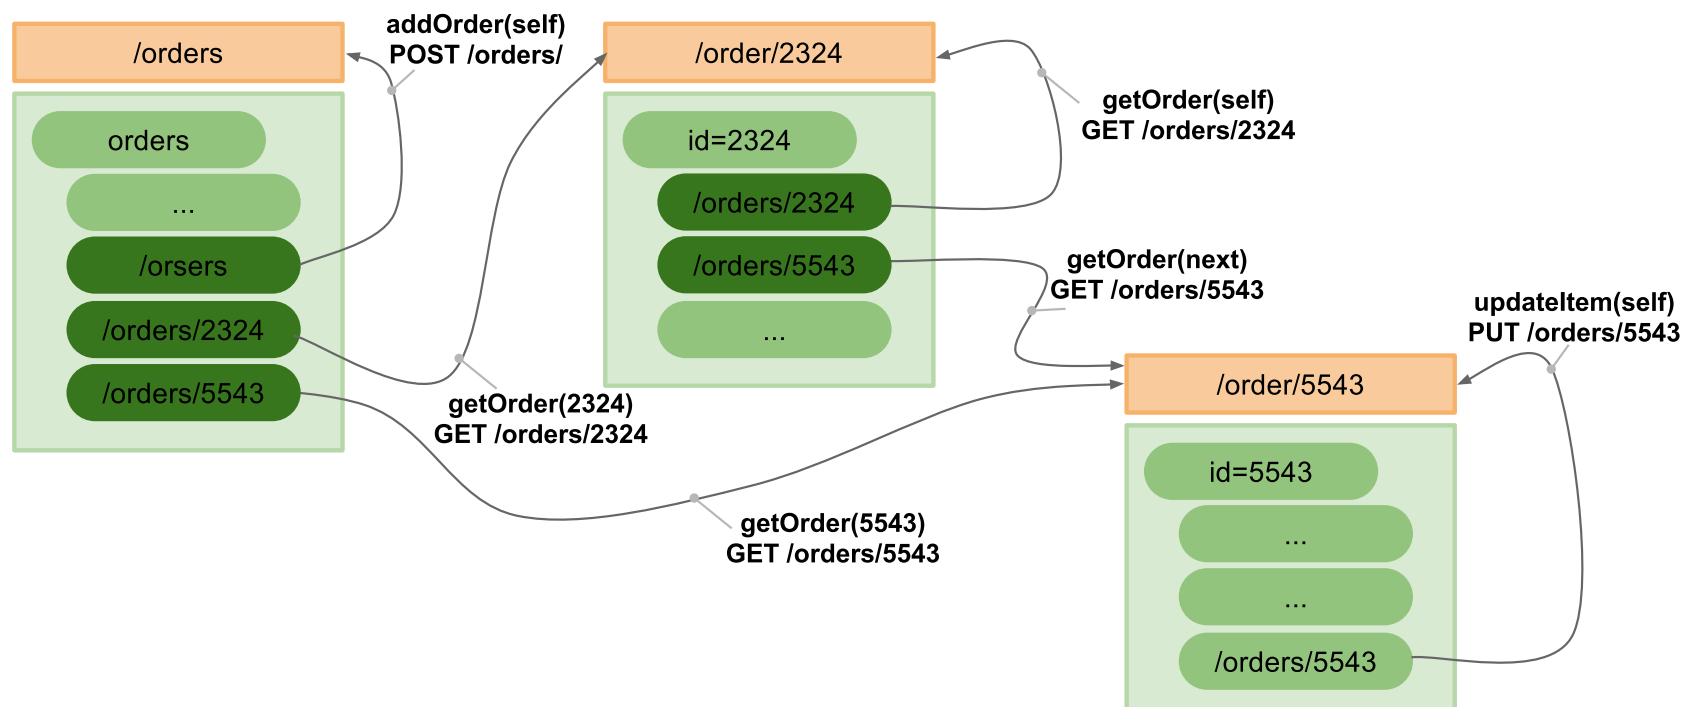
- Persistent Storage
 - *Contains app data*
 - *Data is serialized into resource representation formats*
 - *All sessions may access the data via resource IDs*
 - *Note*
 - *Our simple examples implement a storage in a server memory!*
- Session Memory
 - *Server memory that contains a state of the app*
 - *A session may only access its session memory*
 - *Access through cookies*
 - *Note*
 - *A session memory may be implemented via a persistent storage (such as in Google AppEngine)*

Overview

- HATEOAS
 - *Stateful vs. Stateless*
 - *Links and Preconditions*
- Scalability

Link

- Service operation
 - Applying an access to a link (*GET, PUT, POST, DELETE*)
 - Link: *HTTP method + resource URI + optional link semantics*
- Example: **getOrder**, **addOrder**, and **updateItem**



Atom Links

- Atom Syndication Format
 - *XML-based document format; Atom feeds*
 - *Atom links becoming popular for RESTful applications*
 - *Link structure*
 - rel** – *name of the link*
 - ~ semantics of an operation behind the link*
 - href** – *URI to the resource described by the link*
 - type** – *media type of the resource the link points to*

Link Semantics

- Standard **rel** values
 - *Navigation: next, previous, self*
 - *Does not reflect a HTTP method you can use*
- Extension **rel** values
 - *You can use **rel** to indicate a semantics of an operation*
 - *Example: add item, delete order, update order, etc.*
 - *A client associates this semantics with an operation it may apply at a particular state*
 - *The semantics should be defined by using an URI*

Pagination

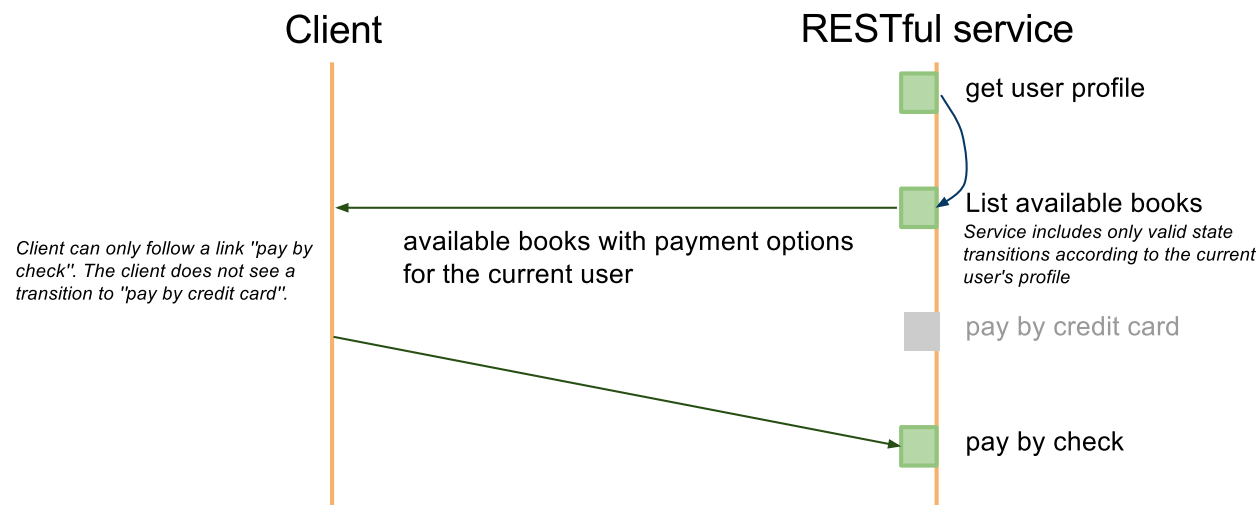
- Dividing a resource into a number of pages
 - *A client retrieves a resource in pages to optimize interactions*
 - *Example: `/orders?page={startPage}&size={numberReturned}`*
 - *A client needs to ask for (or have default values for) a start page and a number of orders to return (must have a pre-defined knowledge)*
- Example `/orders` resource:
 - *client does not need to remember which page of orders it is viewing*

Link Headers

- An alternative to Atom links in resource representations
 - *links defined in HTTP Link header, Web Linking IETF spec* [!\[\]\(849840539e55921a3851a4ff96d7400d_img.jpg\)](#)
 - *They have the same semantics as Atom Links*
 - *Example:*
- Advantages
 - *no need to get the entire document*
 - *no need to parse the document to retrieve links*
 - *use HTTP HEAD only*

Preconditions and HATEOAS

- Precondition
 - Recall Preconditions and effects in MDW
 - A conditions that must hold in a state before an operation can be executed.
- Preconditions in HATEOAS
 - Service in a current state generates only valid transitions that it includes in the representation of the resource.
 - Transition logic is realized at the server-side



Advantages

- Location transparency
 - *only "entry-level" links published to the World*
 - *other links within documents can change without changing client's logic*
 - *HATEOAS may reflect current user's rights in the app*
- Loose coupling
 - *no need for a logic to construct the links*
 - *Clients know to which states they can move via links*

Overview

- HATEOAS
- Scalability
 - *Caching and Revalidation*
 - *Concurrency Control*

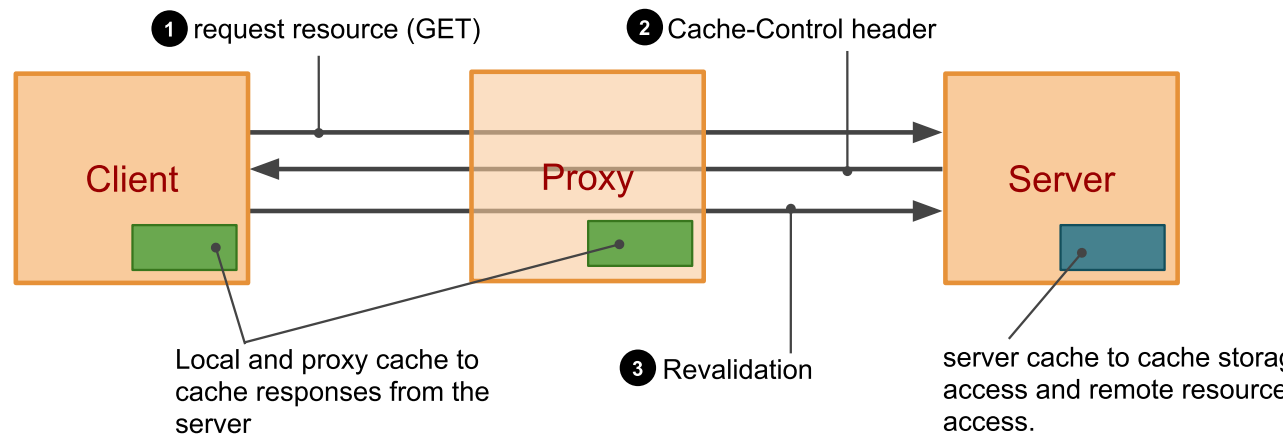
Scalability

- Need for scalability
 - *Huge amount of requests on the Web every day*
 - *Huge amount of data downloaded*
- Some examples
 - *Google, Facebook: 5 billion API calls/day*
 - *Twitter: 3 billions of API calls/day (75% of all the traffic)*
 - *50 million tweets a day*
 - *eBay: 8 billion API calls/month*
 - *Bing: 3 billion API calls/month*
 - *Amazon WS: over 100 billion objects stored in S3*
- Scalability in REST
 - *Caching and revalidation*
 - *Concurrency control*

Overview

- HATEOAS
- Scalability
 - *Caching and Revalidation*
 - *Concurrency Control*

Caching



- Your service should cache:
 - *anytime there is a static resource*
 - *even there is a dynamic resource*
 - *with chances it updates often*
 - *you can force clients to always revalidate*
- three steps:
 - *client GETs the resource representation*
 - *server controls how it should cache through **Cache-Control** header*
 - *client revalidates the content via conditional GET*

Cache Headers

- **Cache-Control** response header
 - *controls over local and proxy caches*
 - **private** – *no proxy should cache, only clients can*
 - **public** – *any intermediary can cache (proxies and clients)*
 - **no-cache** – *the response should not be cached. If it is cached, the content should always be revalidated.*
 - **no-store** – *can cache but should not store persistently. When a client restarts, content is lost*
 - **no-transform** – *no transformation of cached data; e.g. compressions*
 - **max-age**, **s-maxage** *a time in seconds how long the cache is valid; s-maxage for proxies*
- **Last-Modified** and **ETag** response headers
 - *Content last modified date and a content entity tag*
- **If-Modified-Since** and **If-None-Match** request headers
 - *Content revalidation (conditional GET)*

Example Date Revalidation

- Cache control example:
 - *only client can cache, must not be stored on the disk, the cache is valid for 200 seconds.*
- Revalidation (conditional GET) example:
 - *A client revalidates the cache after 200 seconds.*

Entity Tags

- Signature of the response body
 - *A hash such as MD5*
 - *A sequence number that changes with any modification of the content*
- Types of tag
 - *Strong ETag: reflects the content bit by bit*
 - *Weak ETag: reflects the content "semantically"*
 - *The app defines the meaning of its weak tags*
- Example content revalidation with ETag

Design Suggestions

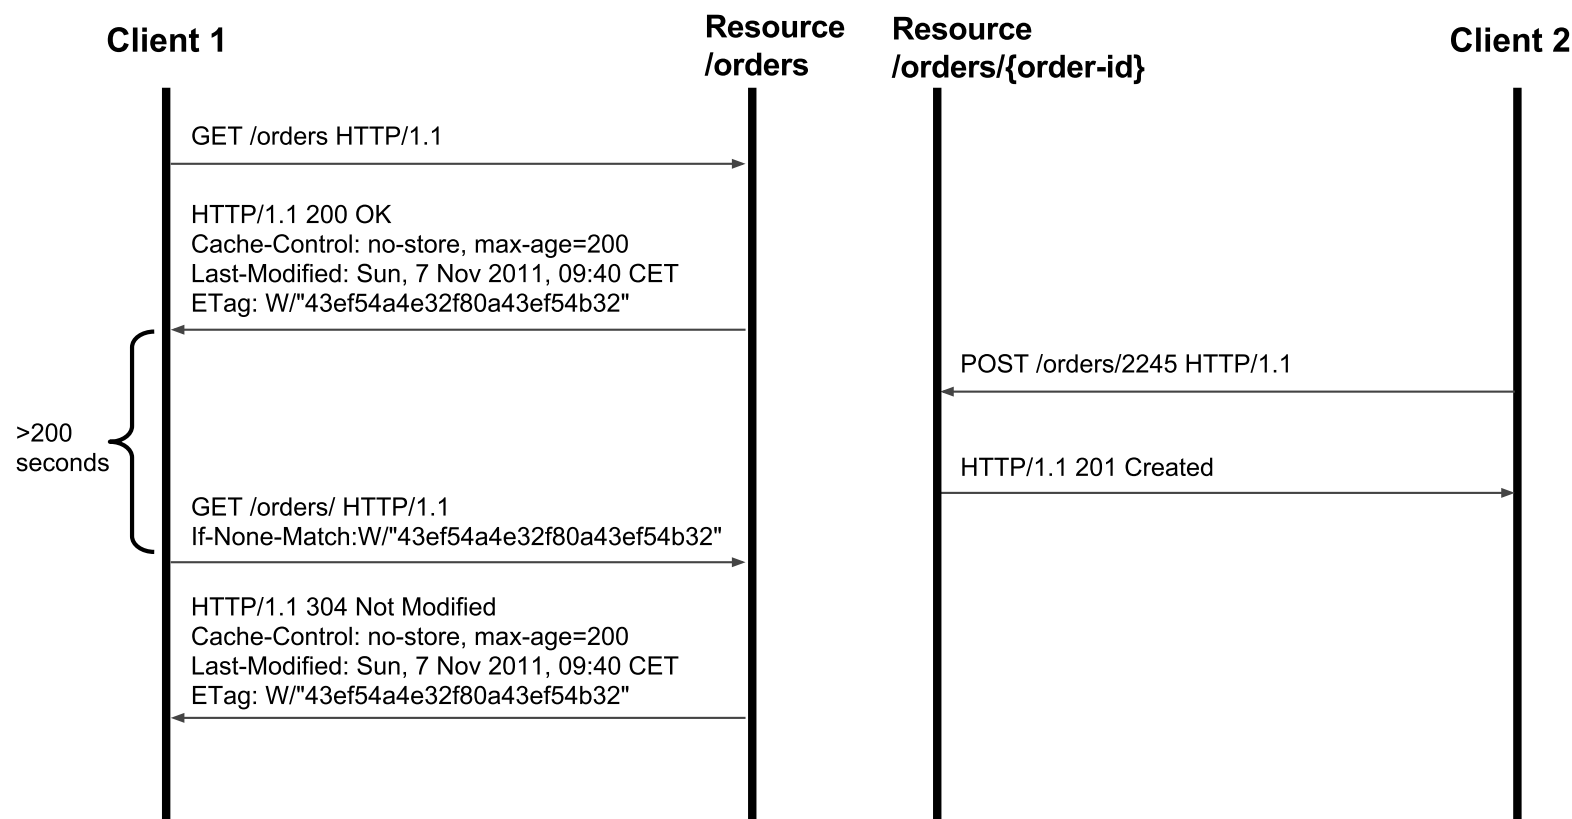
- Composed resources use weak ETags
 - For example */orders*
 - a composed resource that contains a summary information
 - changes to an order's items will not change semantics of */orders*
 - It is usually not possible to perform updates on these resources
- Non-composed resources use strong ETags
 - For example */orders/{order-id}*
 - They can be updated
- Further notes
 - Server should send both *Last-Modified* and *ETag* headers
 - If client sends both *If-Modified-Since* and *If-None-Match*, *ETag* validation takes preference

Weak ETag Example

- App specific, */orders* resource example
- Weak ETag compute function example
 - *Any modification to an order's items is not significant for /orders:*

Weak ETag Revalidation

- Updating **/orders** resource
 - **POST** **/orders/{order-id}** *inserts a new item to an order*
 - *Any changes to orders' items will not change the Weak ETag*



Overview

- HATEOAS
- Scalability
 - *Caching and Revalidation*
 - *Concurrency Control*

Concurrency

- Two clients may update the same resource

1) a client GETs a resource GET /orders/5545

2) the client modifies the resource

3) the client updates the resource via PUT /orders/5545 HTTP/1.1

What happens if another client updates the resource between 1) and 3) ?

- Concurrency control

- Conditional PUT

- Update the resource only if it has not changed since a specified date or a specified ETag matches the resource content

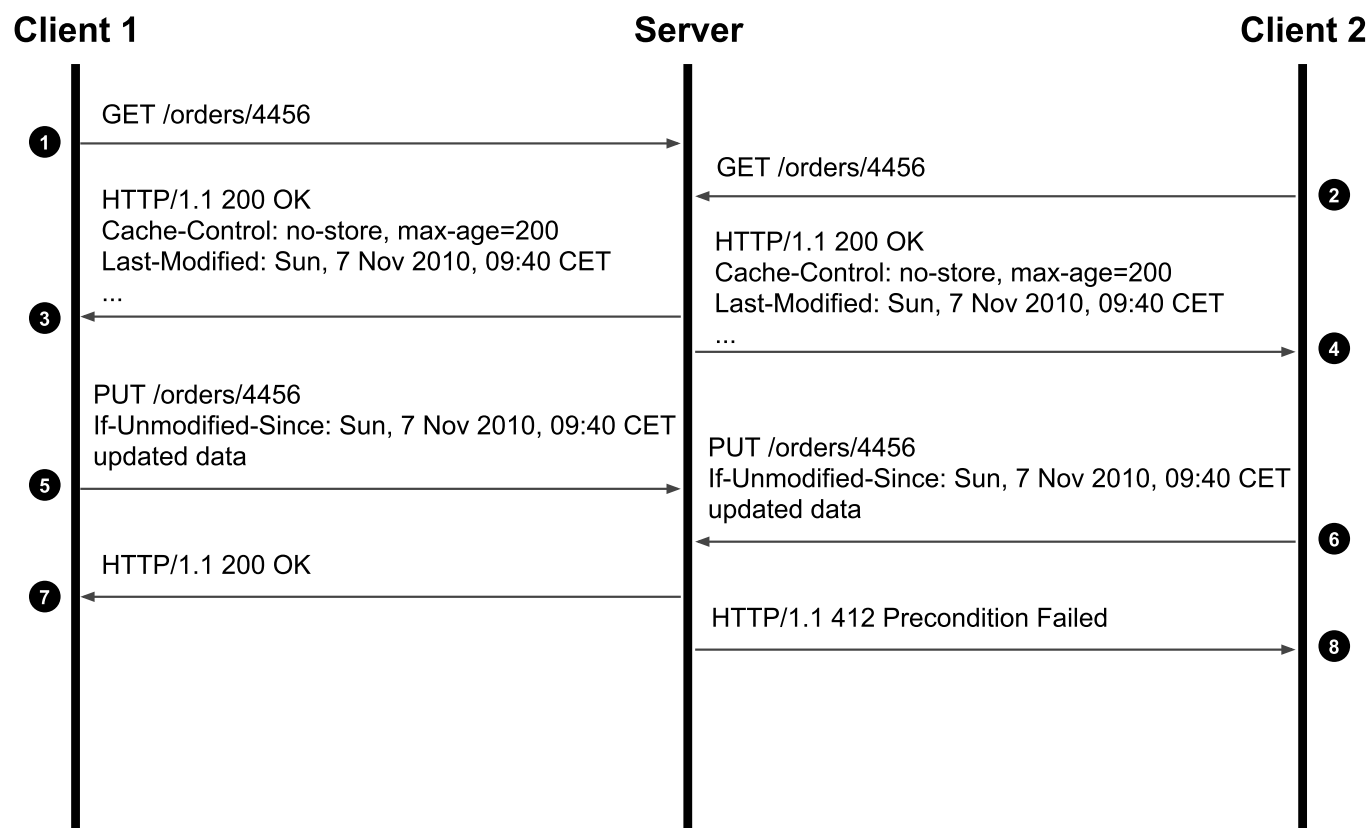
- If-Unmodified-Since and If-Match headers

- Response to conditional PUT:

- 200 OK if the PUT was successful

- 412 Precondition Failed if the resource was updated in the meantime.

Concurrency Control Protocol



- Conditional PUT and ETags
 - *Conditional PUT must always use strong entity tags or date validation*

Overview

- HATEOAS
- Scalability
 - *Documentation*

Documentation

- RESTful API Documentation
 - *Until recently, not a standard way, only good practices*
 - *and only textual, not in a formal language*
 - *there were attempts such as WADL, hREST*
 - *it is even possible to use WSDL 2.0*
 - *Today, Swagger and Open API Specification*
- Client libraries in major languages
 - *JavaScript, Java, ...*
 - *these could be documented*
 - *they hide protocol details*
- Best practices in RESTful API documentation
 - *learn from Google, Twitter, and others*

Best Practices

- Include resource diagram
 - *in UML, with links*
- For each resource, describe
 - *URI with parameters, such as*
`http://company.com/orders/{order-id}`
 - *definition of the parameters*
 - *list of properties (attributes), with values, link to XML Schema*
 - *representations you support (XML, JSON)*
 - *sample request*
 - *sample response in representations you support*
 - *error codes*
- Make sure
 - *people can copy sample code and run it in a browser or by using*
`curl`

Swagger Overview

- Emerging standard
 - *Started as a private company effort (SmartBear)*
 - *Recently became so popular and evolved to a community effort*
 - *Open API Specification under Apache Foundation*
 - *Google, IBM, 3Scale, ...*
- Guiding Principles
 - *A minimal effort to describe an API*
 - *API description should be generated, e.g. via code annotations*
 - *It can always be written manually too*
 - *A minimal effort to write clients*
 - *Sanbox comes out-of-the-box*

Swagger API Description

- Server
 - Server provides a **Resource Listing** at `/api-docs`
 - For each resource, there is an **API Declaration**
- Resource Listing
 - JSON Representation

Swagger API Description

- API Declaration
 - *JSON Representation*