

# Web 2.0

## Lecture 4: HATEOAS, Scalability and Description

**doc. Ing. Tomáš Vitvar, Ph.D.**

tomas@vitvar.com • @TomasVitvar • <http://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <http://vitvar.com/courses/w20>



Evropský sociální fond  
Praha & EU: Investujeme do vaší budoucnosti

Modified: Mon Mar 20 2017, 20:22:04  
Humla v0.3

# REST Core Principles

- REST architectural style defines constraints
  - *if you follow them, they help you to achieve a good design, interoperability and scalability.*
- Constraints
  - *Client/Server*
  - *Statelessness*
  - *Cacheability*
  - *Layered system*
  - *Uniform interface*
- Guiding principles
  - *Identification of resources*
  - *Representations of resources and self-descriptive messages*
  - *Hypermedia as the engine of application state (HATEOAS)*

# Overview

- HATEOAS
  - *Stateful vs. Stateless*
  - *Links and Preconditions*
- Scalability

# HATEOAS

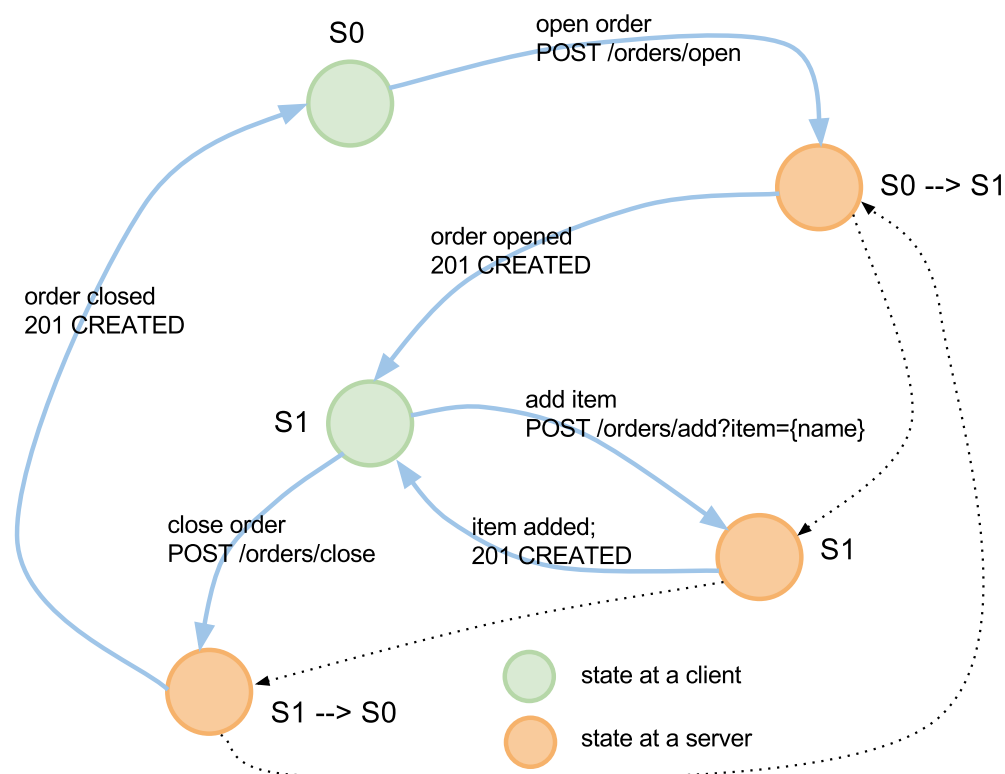
- HATEOAS = Hypertext as the Engine for Application State
  - *The REST core principle*
  - **Hypertext**
    - *Hypertext is a representation of a resource with **links***
    - *A link is an URI of a resource*
    - *Applying an access to a resource via its link = state transition*
- Statelessness
  - *A service does not use a memory to remember a state*
  - *HATEOAS enables stateless implementation of services*

# Overview

- HATEOAS
  - *Stateful vs. Stateless*
  - *Links and Preconditions*
- Scalability

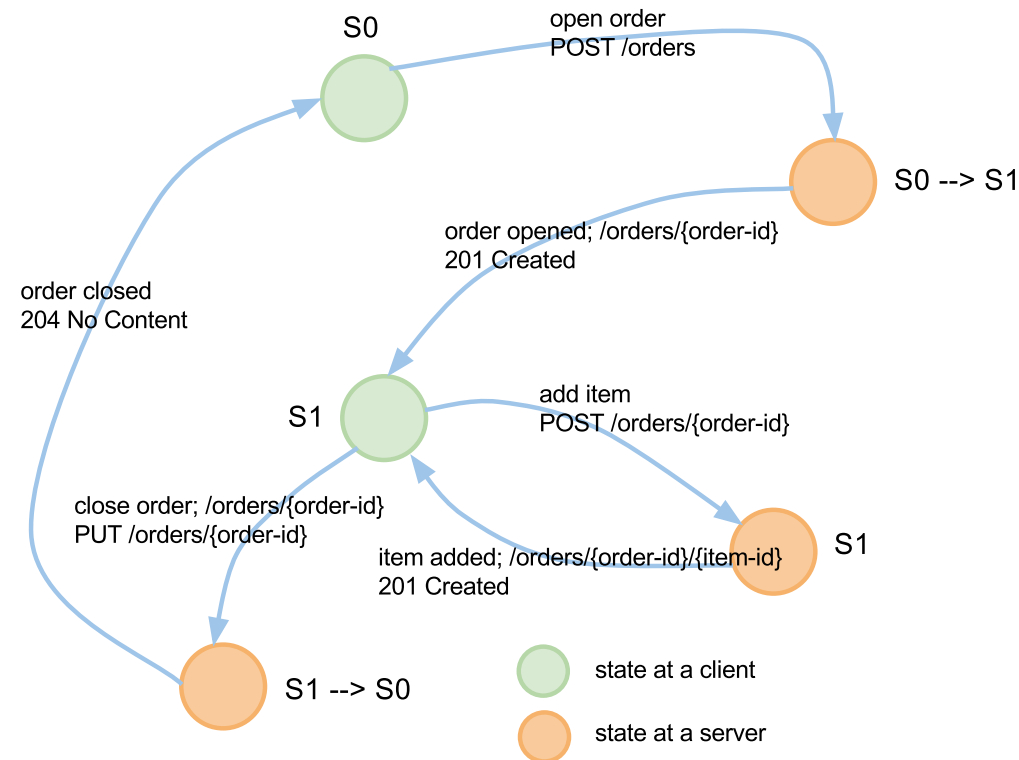
# Stateful server

- Sessions to store the application state
  - Recall HTTP state management in MDW
  - The app uses a server memory to remember the state
  - when server restarts, the app state is lost



# Stateless server

- HTTP and hypermedia to transfer the app state
  - *Does not use a server memory to remember the app state*
  - *State transferred between a client and a service via HTTP metadata and resources' representations*



# Persistent Storage and Session Memory

- Persistent Storage
  - *Contains app data*
  - *Data is serialized into resource representation formats*
  - *All sessions may access the data via resource IDs*
  - *Note*
    - *Our simple examples implement a storage in a server memory!*
- Session Memory
  - *Server memory that contains a state of the app*
  - *A session may only access its session memory*
  - *Access through cookies*
  - *Note*
    - *A session memory may be implemented via a persistent storage (such as in Google AppEngine)*

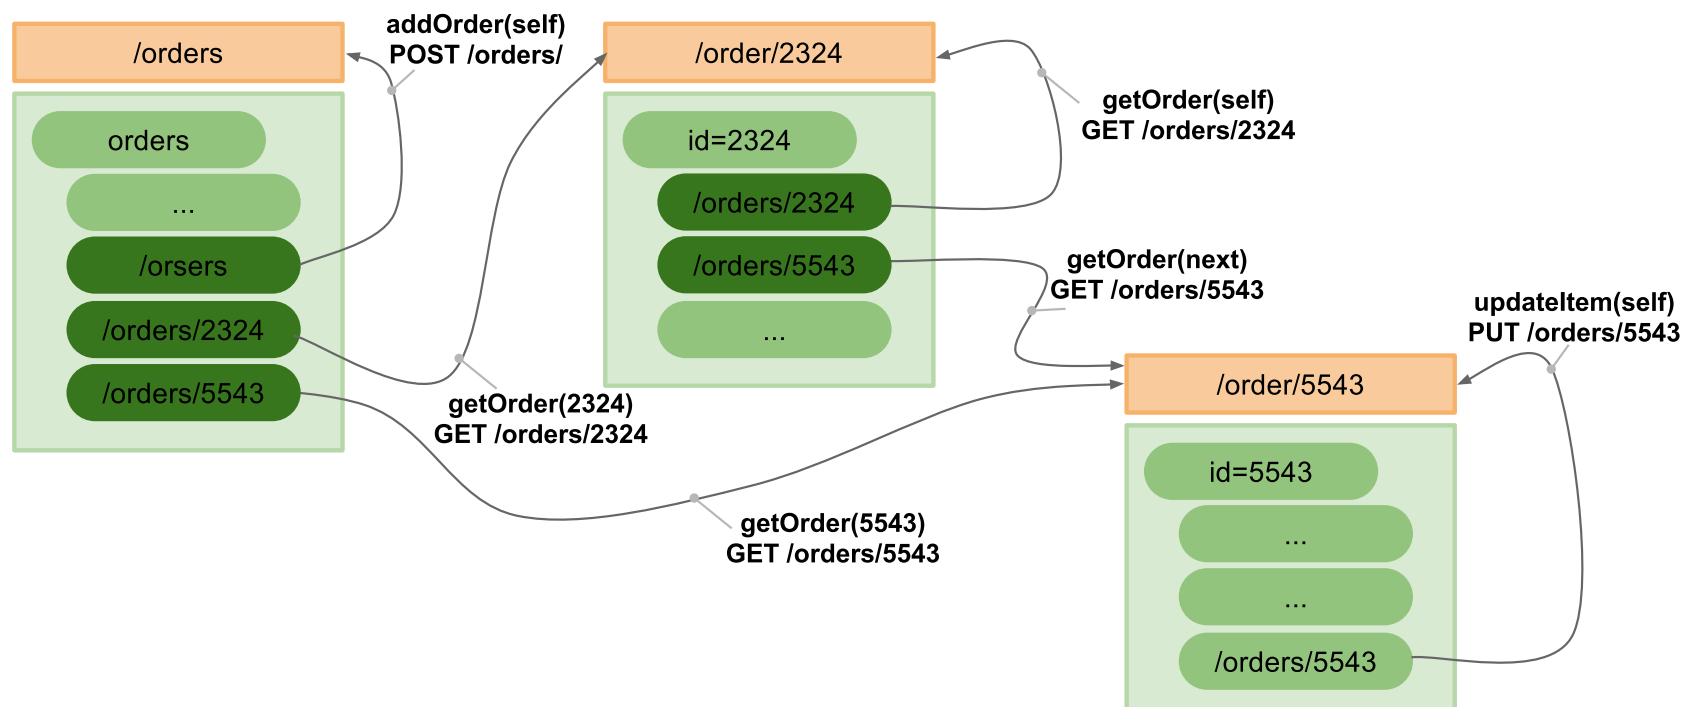


# Overview

- HATEOAS
  - *Stateful vs. Stateless*
  - *Links and Preconditions*
- Scalability

# Link

- Service operation
  - Applying an access to a link (*GET, PUT, POST, DELETE*)
  - Link: *HTTP method + resource URI + optional link semantics*
- Example: **getOrder**, **addOrder**, and **updateItem**



# Atom Links

- Atom Syndication Format
  - *XML-based document format; Atom feeds*
  - *Atom links becoming popular for RESTful applications*
  - *Link structure*
    - rel** – *name of the link*
    - ~ semantics of an operation behind the link*
    - href** – *URI to the resource described by the link*
    - type** – *media type of the resource the link points to*

# Link Semantics

- Standard **rel** values
  - *Navigation: next, previous, self*
  - *Does not reflect a HTTP method you can use*
- Extension **rel** values
  - *You can use **rel** to indicate a semantics of an operation*
  - *Example: add item, delete order, update order, etc.*
  - *A client associates this semantics with an operation it may apply at a particular state*
  - *The semantics should be defined by using an URI*

# Pagination

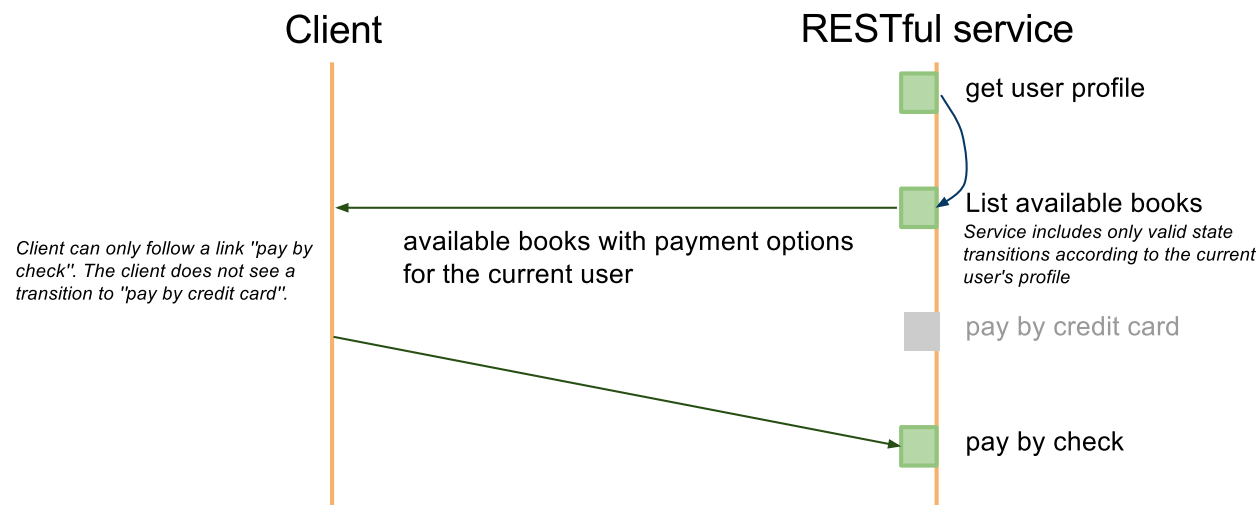
- Dividing a resource into a number of pages
  - *A client retrieves a resource in pages to optimize interactions*
  - *Example: `/orders?page={startPage}&size={numberReturned}`*
  - *A client needs to ask for (or have default values for) a start page and a number of orders to return (must have a pre-defined knowledge)*
- Example `/orders` resource:
  - *client does not need to remember which page of orders it is viewing*

# Link Headers

- An alternative to Atom links in resource representations
  - *links defined in HTTP Link header, Web Linking IETF spec* [!\[\]\(849840539e55921a3851a4ff96d7400d\_img.jpg\)](#)
  - *They have the same semantics as Atom Links*
  - *Example:*
- Advantages
  - *no need to get the entire document*
  - *no need to parse the document to retrieve links*
  - *use HTTP HEAD only*

# Preconditions and HATEOAS

- Precondition
  - Recall Preconditions and effects in MDW
    - A conditions that must hold in a state before an operation can be executed.
- Preconditions in HATEOAS
  - Service in a current state generates only valid transitions that it includes in the representation of the resource.
  - Transition logic is realized at the server-side



# Advantages

- Location transparency
  - *only "entry-level" links published to the World*
  - *other links within documents can change without changing client's logic*
  - *HATEOAS may reflect current user's rights in the app*
- Loose coupling
  - *no need for a logic to construct the links*
  - *Clients know to which states they can move via links*



# Overview

- HATEOAS
- Scalability
  - *Caching and Revalidation*
  - *Concurrency Control*

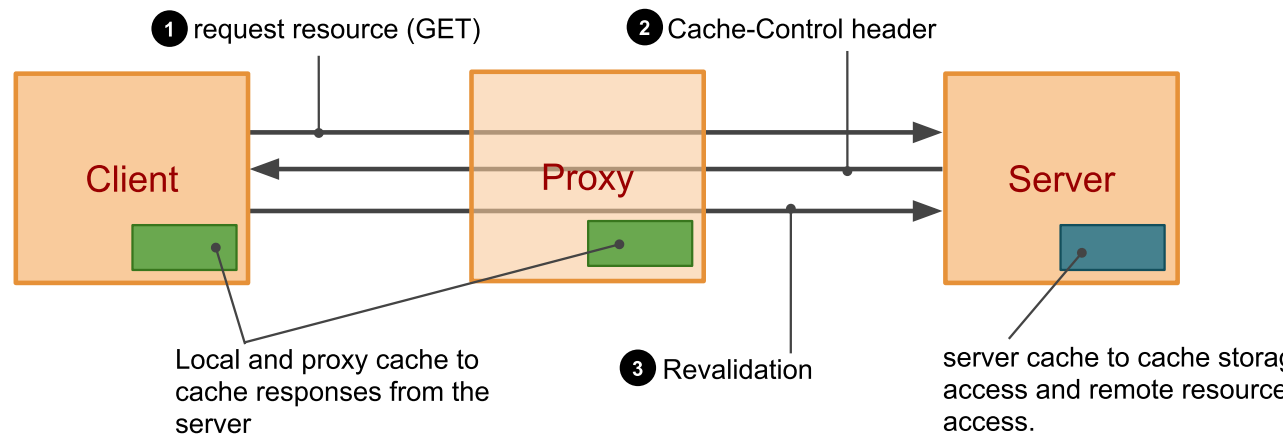
# Scalability

- Need for scalability
  - *Huge amount of requests on the Web every day*
  - *Huge amount of data downloaded*
- Some examples
  - *Google, Facebook: 5 billion API calls/day*
  - *Twitter: 3 billions of API calls/day (75% of all the traffic)*
    - *50 million tweets a day*
  - *eBay: 8 billion API calls/month*
  - *Bing: 3 billion API calls/month*
  - *Amazon WS: over 100 billion objects stored in S3*
- Scalability in REST
  - *Caching and revalidation*
  - *Concurrency control*

# Overview

- HATEOAS
- Scalability
  - *Caching and Revalidation*
  - *Concurrency Control*

# Caching



- Your service should cache:
  - *anytime there is a static resource*
  - *even there is a dynamic resource*
    - *with chances it updates often*
    - *you can force clients to always revalidate*
- three steps:
  - *client GETs the resource representation*
  - *server controls how it should cache through **Cache-Control** header*
  - *client revalidates the content via conditional GET*

# Cache Headers

- **Cache-Control** response header
  - *controls over local and proxy caches*
  - **private** – *no proxy should cache, only clients can*
  - **public** – *any intermediary can cache (proxies and clients)*
  - **no-cache** – *the response should not be cached. If it is cached, the content should always be revalidated.*
  - **no-store** – *can cache but should not store persistently. When a client restarts, content is lost*
  - **no-transform** – *no transformation of cached data; e.g. compressions*
  - **max-age**, **s-maxage** *a time in seconds how long the cache is valid; s-maxage for proxies*
- **Last-Modified** and **ETag** response headers
  - *Content last modified date and a content entity tag*
- **If-Modified-Since** and **If-None-Match** request headers
  - *Content revalidation (conditional GET)*

# Example Date Revalidation

- Cache control example:
  - *only client can cache, must not be stored on the disk, the cache is valid for 200 seconds.*
- Revalidation (conditional GET) example:
  - *A client revalidates the cache after 200 seconds.*

# Entity Tags

- Signature of the response body
  - *A hash such as MD5*
  - *A sequence number that changes with any modification of the content*
- Types of tag
  - *Strong ETag: reflects the content bit by bit*
  - *Weak ETag: reflects the content "semantically"*
    - *The app defines the meaning of its weak tags*
- Example content revalidation with ETag

# Design Suggestions

- Composed resources use weak ETags
  - For example */orders*
    - a composed resource that contains a summary information
    - changes to an order's items will not change semantics of */orders*
  - It is usually not possible to perform updates on these resources
- Non-composed resources use strong ETags
  - For example */orders/{order-id}*
  - They can be updated
- Further notes
  - Server should send both *Last-Modified* and *ETag* headers
  - If client sends both *If-Modified-Since* and *If-None-Match*, *ETag* validation takes preference

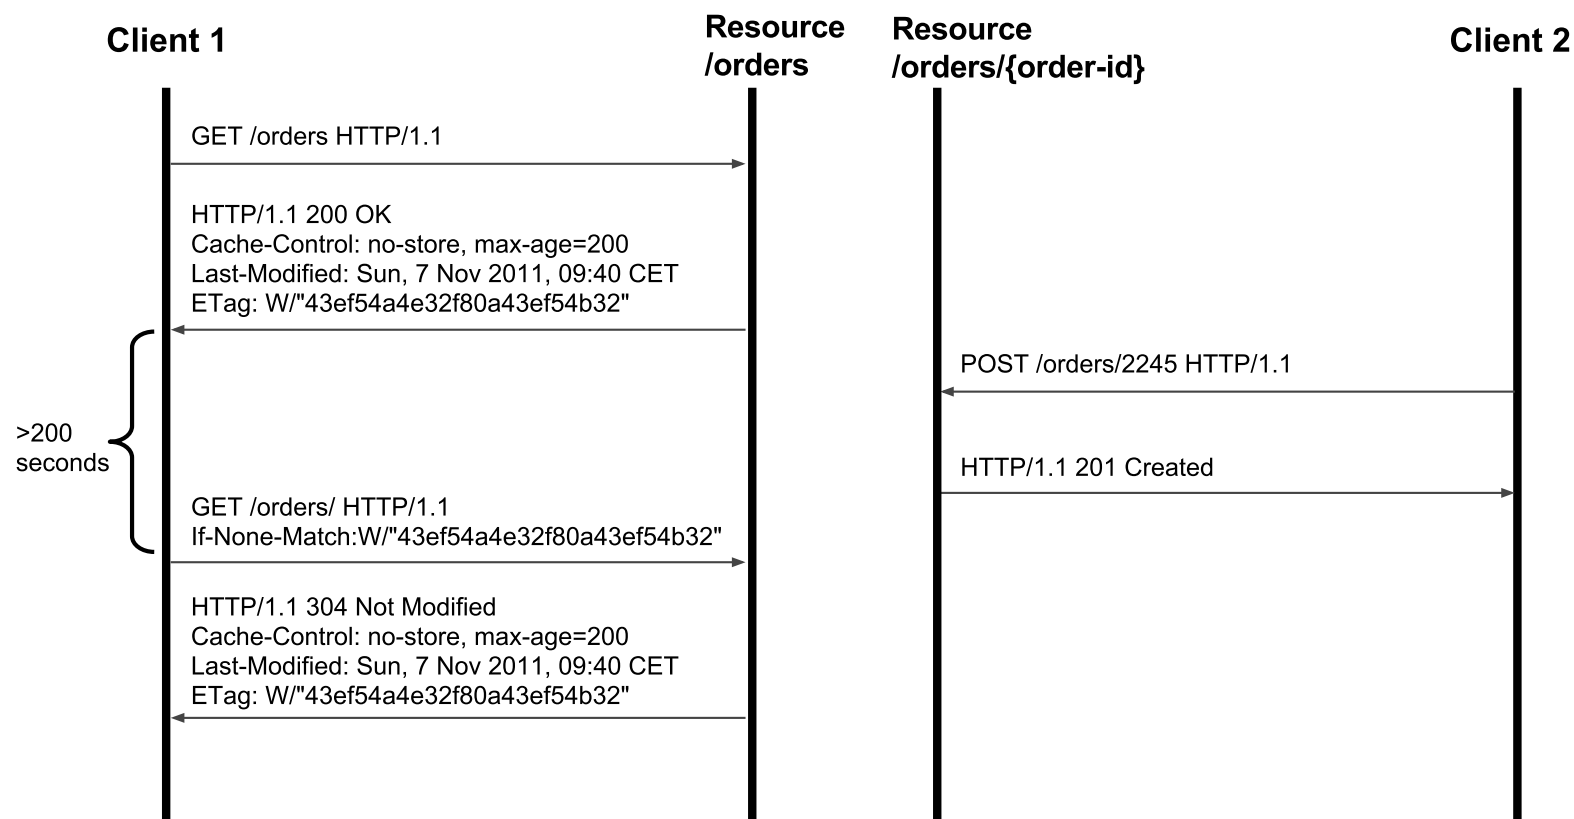


# Weak ETag Example

- App specific, */orders* resource example
- Weak ETag compute function example
  - *Any modification to an order's items is not significant for /orders:*

# Weak ETag Revalidation

- Updating **/orders** resource
  - **POST** **/orders/{order-id}** *inserts a new item to an order*
  - *Any changes to orders' items will not change the Weak ETag*



# Overview

- HATEOAS
- Scalability
  - *Caching and Revalidation*
  - *Concurrency Control*

# Concurrency

- Two clients may update the same resource

*1) a client GETs a resource GET /orders/5545*

*2) the client modifies the resource*

*3) the client updates the resource via PUT /orders/5545 HTTP/1.1*

*What happens if another client updates the resource between 1) and 3) ?*

- Concurrency control

- Conditional PUT

- Update the resource only if it has not changed since a specified date or a specified ETag matches the resource content

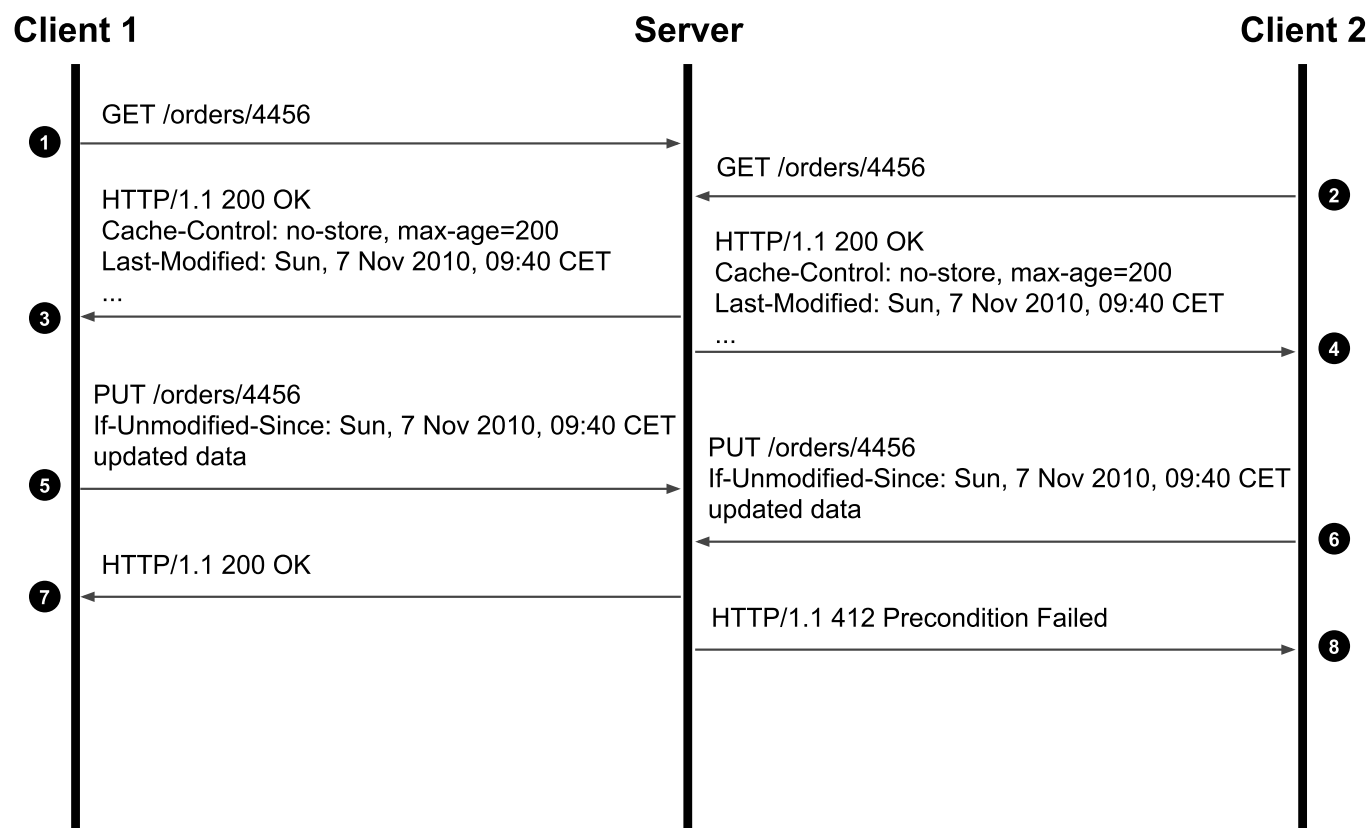
- If-Unmodified-Since and If-Match headers

- Response to conditional PUT:

- 200 OK if the PUT was successful

- 412 Precondition Failed if the resource was updated in the meantime.

# Concurrency Control Protocol



- Conditional PUT and ETags
  - *Conditional PUT must always use strong entity tags or date validation*

# Overview

- HATEOAS
- Scalability
  - *Documentation*

# Documentation

- RESTful API Documentation
  - *Until recently, not a standard way, only good practices*
  - *and only textual, not in a formal language*
    - *there were attempts such as WADL, hREST*
    - *it is even possible to use WSDL 2.0*
  - *Today, Swagger and Open API Specification*
- Client libraries in major languages
  - *JavaScript, Java, ...*
  - *these could be documented*
  - *they hide protocol details*
- Best practices in RESTful API documentation
  - *learn from Google, Twitter, and others*

# Best Practices

- Include resource diagram
  - *in UML, with links*
- For each resource, describe
  - *URI with parameters, such as*  
`http://company.com/orders/{order-id}`
  - *definition of the parameters*
  - *list of properties (attributes), with values, link to XML Schema*
  - *representations you support (XML, JSON)*
  - *sample request*
  - *sample response in representations you support*
  - *error codes*
- Make sure
  - *people can copy sample code and run it in a browser or by using*  
`curl`



# Swagger Overview

- Emerging standard
  - *Started as a private company effort (SmartBear)*
  - *Recently became so popular and evolved to a community effort*
    - *Open API Specification under Apache Foundation*
    - *Google, IBM, 3Scale, ...*
- Guiding Principles
  - *A minimal effort to describe an API*
    - *API description should be generated, e.g. via code annotations*
    - *It can always be written manually too*
  - *A minimal effort to write clients*
  - *Sanbox comes out-of-the-box*

# Swagger API Description

- Server
  - Server provides a **Resource Listing** at `/api-docs`
  - For each resource, there is an **API Declaration**
- Resource Listing
  - JSON Representation

# Swagger API Description

- API Declaration
  - *JSON Representation*