Web 2.0

Lecture 8: Protocols for the Realtime Web

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • http://vitvar.com

Czech Technical University in Prague

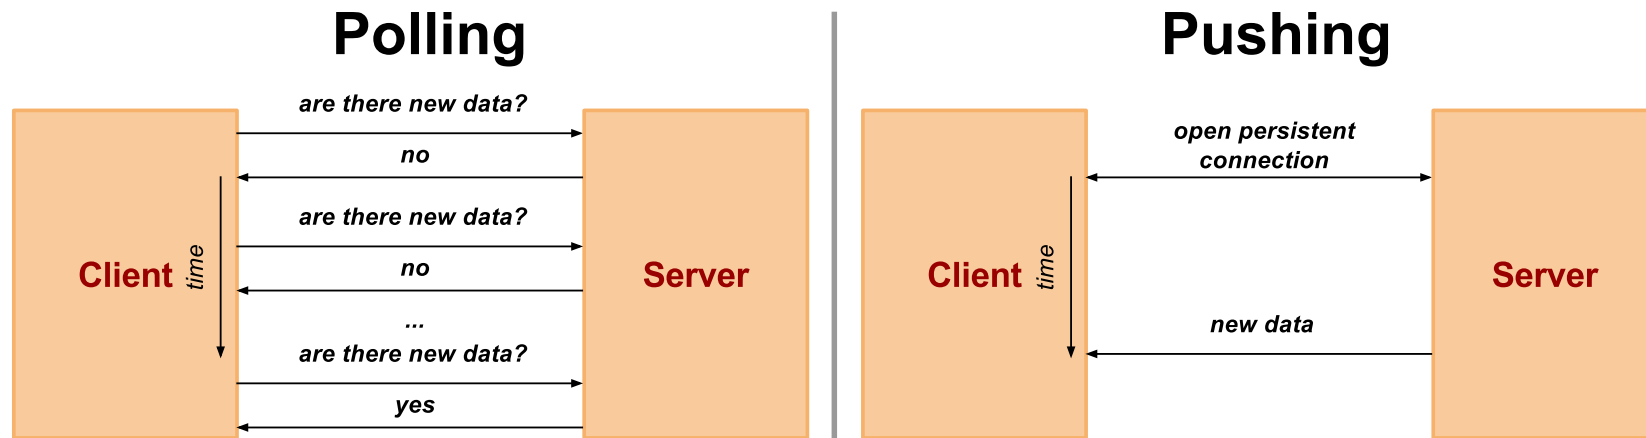Faculty of Information Technologies • Software and Web Engineering • http://vitvar.com/courses/w20

Overview

- Long-polling and Streaming
- WebSocket Protocol
- New I/O Model

Pushing and Polling

Polling

```
are there new data?
                          →
         no
                          ←
are there new data?
                          →
         no
                          ←
        ...
are there new data?
                          →
        yes
                          ←
```

Client *time* — **Server**

Pushing

```
open persistent
   connection
                          ↔

    new data
                          ←
```

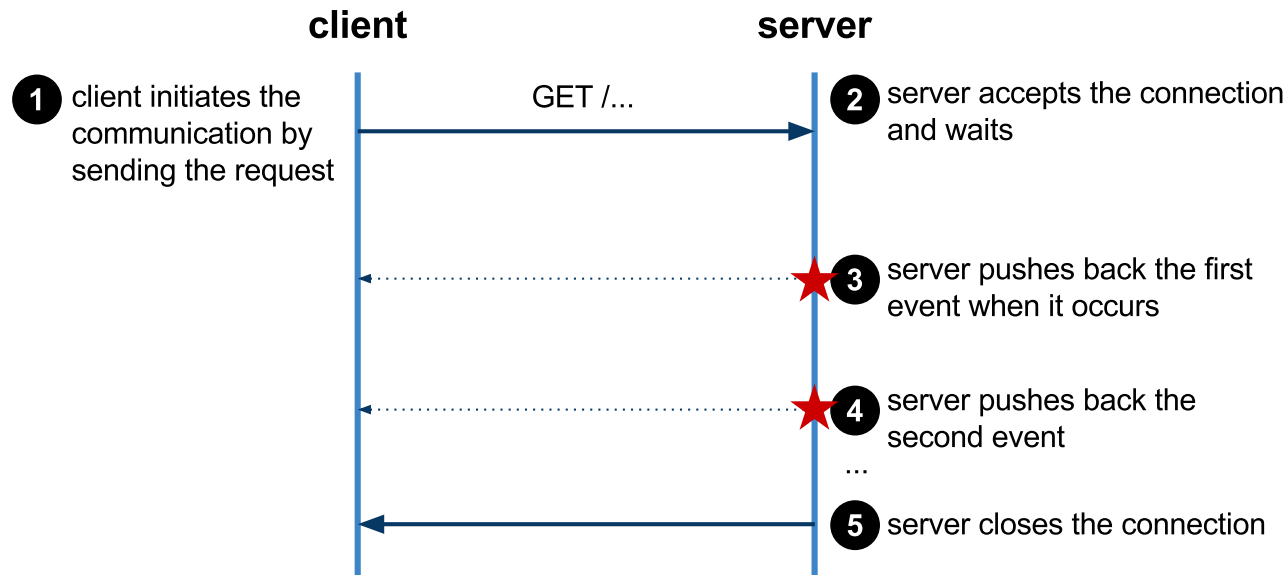**Client** *time* — **Server**

- Conceptual basis in messaging architectures
  - *event-driven architectures (EDA)*
- **HTTP is a request-response protocol**
  - *response cannot be sent without request*
  - *server cannot initiate the communication*
- **Polling** – client periodically checks for updates on the server
- **Pushing** – updates from the server (also called COMET)
  = **long polling** – server holds the request for some time
  = **streaming** – server sends updates without closing the socket

HTTP Long Polling


```
            client                              server
              |                                   |
              |          GET /...                 |
client initiates a communication                 | server does not respond immediately
by sending a request  |─────────────────────────>|  but waits until a new event occurs.
              |                                   |
              |                                   | when a new event occurs, server
              |<─────────────────────────★       |  sends a response message and
              |                                   |  closes the connection.
client resends the request, it   GET /...         |
may wait 2-3 seconds to save  |──────────────────>|
resources     |                                   |
              |                                   |
              |<─────────────────────────★       |
              |              ...                  |
```
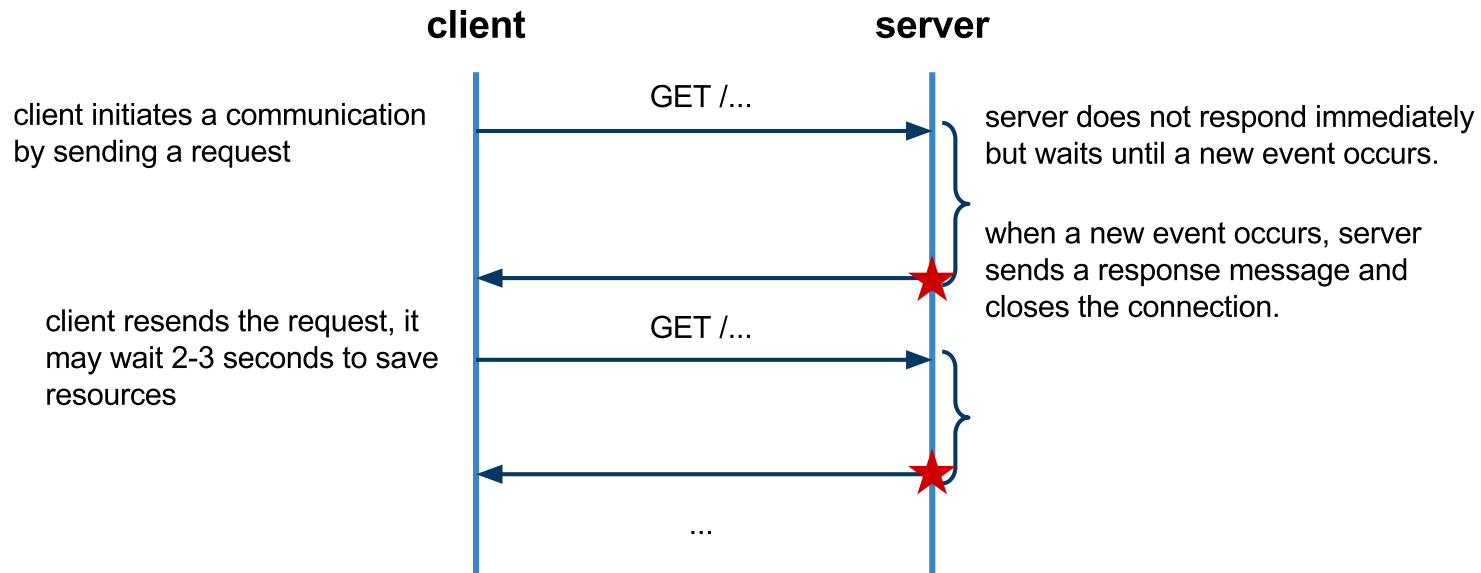
- Server holds long-poll requests
  - *server responds when an event or a timeout occurs*
  - *saves computing resources at the server as well as network resources*
  - *can be applied over HTTP persistent and non-persistent communication*

- Issues:
  - *maximum time of the request processing at the server*
  - *concurrent requests processing at the server*

HTTP Streaming

- server deffers the response until an event or timeout is available
- when an event is available, server sends it back to client as part of the response; this does not terminate the connection
- server is able to send pieces of response w/o terminating the conn.
  - *using* `transfer-encoding` *header in HTTP 1.1*
  - *using End of File in HTTP 1.0
    (server omits* `content-lenght` *in the response)*

Chunked Response

- Transfer encoding `chunked`
  - *It allows to send multiple sets of data over a single connection*
  - *a chunk represents data for the event*
  - *Each chunk starts with hexadecimal value for length*
  - *End of response is marked with the chunk length of 0*

- Steps:
  - *server sends HTTP headers and the first chunk (step 3)*
  - *server sends second and subsequent chunk of data (step 4)*
  - *server terminates the connection (step 5)*

Issues with Chunked Response

- ## Chunks vs. Events
  - *chunks cannot be considered as app messages (events)*
  - *intermediaries might "re-chunk" the message stream*
    → e.g., combining different chunks into a longer one

- ## Client Buffering
  - *clients may buffer all data chunks before they make the response available to the client application*

- ## HTTP streaming in browsers
  - *Server-sent events*

Server-Sent Events

- W3C specification
  - *part of HTML5 specs, see Server-Sent Events* ⬈
  - *API to handle HTTP streaming in browsers by using DOM events*
  - *transparent to underlying HTTP streaming mechanism*
    - → *can use both chunked messages and EOF*
  - *same origin policy applies*

- `EventSource` interface
  - *event handlers:* `onopen`, `onmessage`, `onerror`
  - *constructor* `EventSource(url)` – *creates and opens the stream*
  - *method* `close()` – *closes the connection*
  - *attribute* `readyState`
    - → `CONNECTING` – *The connection has not yet been established, or it was closed and the user agent is reconnecting.*
    - → `OPEN` – *The user agent has an open connection and is dispatching events as it receives them.*
    - → `CLOSED` – *The conn. is not open, the user agent is not reconnecting.*

Example

- Initiating `EventSource`

- Defining event handlers
  - *when the conn. is closed, the browser reconnects every ~3 seconds*
    → *can be changed using* `retry` *attribute in the message data*

Event Stream Format

- Format
  - *response's* `content-type` *must be* `text/event-stream`
  - *every line starts with* `data:`*, event message terminates with 2* `\n` *chars.*
  - *every message may have associated* `id` *(is optional)*

- JSON data in multiple lines of the message

- Changing the reconnection time
  - *default is 3 seconds*

Server-side implementation

- Java Servlet
  - *method* `doGet`

Other Technologies

- Cross-document messaging

script in document A

```
var o document.getElementById("ifr");
o.contentWindow.postMessage("Hello world",
    "http://example.org/")
```

document A
`<iframe id="ifr" ...>`
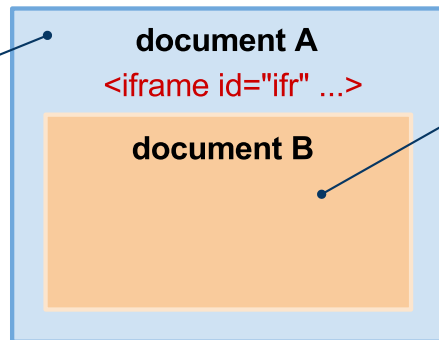
document B

script in document B

```
window.addEventListener('message', receiver, false);
function receiver(e) {
  if (e.origin == 'http://example.com') {
    if (e.data == 'Hello world') {
      e.source.postMessage('Hello', e.origin);
    } else {
      alert(e.data);
    }
  }
}
```

- *The use of Cross Document Messaging for streaming*
  1. *The client loads a streaming resource in a hidden* `iframe`
  2. *The server pushes a JavaScript code to the* `iframe`
  3. *The browser executes the code as it arrives from the server*
  4. *The embedded iframe's code posts a message to the upper document*

- Channel API
  - *Google Technology for streaming API for AppEngine*
  - *not based on HTTP streaming*
  - *utilizes XMPP capabilities + hidden iframe at client-side*

Overview

- Long-polling and Streaming
- WebSocket Protocol
- New I/O Model

WebSocket

- Specifications
  - *IETF defines WebSocket Protocol* ⬈
  - *W3C defines WebSocket API* ⬈

- Design principles
  - *a new protocol*
    - *→ browsers, web servers, and proxy servers need to support it*
  - *a layer on top of TCP*
  - *bi-directional communication between client and servers*
    - *→ low-latency apps without HTTP overhead*
  - *Web origin-based security model for browsers*
    - *→ same origin policy, cross-origin resource sharing*
  - *support multiple server-side endpoints*

- Two phases
  - *Handshake – as an **upgrade** of a HTTP connection*
  - *data transfer – the protocol-specific on-the-wire data transfer*

Handshake – Request

- Request
  - *client sends a following HTTP request to upgrade the connection to WebSocket*
  - `Connection` – *request to upgrade the protocol*
  - `Upgrade` – *protocol to upgrade to*
  - `Sec-WebSocket-Key` – *a client key for later validation*
  - `Sec-WebSocket-Origin` – *origin of the request*
  - `Sec-WebSocket-Protocol` – *list of sub-protocols that client supports (proprietary)*

Handshake – Response

- Response
  - *server accepts the request and responds as follows*
    - → `101 Switching Protocols` – *status code for a successful upgrade*
    - → `Sec-WebSocket-Protocol` – *a sub-protocol that the server selected from the list of protocols in the request*
    - → `Sec-WebSocket-Accept` – *a key to prove it has received a client WebSocket handshake request*
  - *Formula to compute* `Sec-WebSocket-Accept`
    - → `SHA-1` – *hashing function*
    - → `Base64Encode` – *Base64 encoding function*
    - → `"258EAFA5-E914-47DA-95CA-C5AB0DC85B11"` – *magic number*

Data Transfer

- After successful handshake
  - *socket between the client and the "resource" at the server is established*
  - *client and the server can both read and write from/to the socket*
  - *No HTTP headers overhead*

- Data Framing
  - *defines a format for data transmitted in TCP packets*
  - *payload length, closing frame, ping, pong, type of data (text/binary), etc. and payload (message data)*

WebSocket API

- Client-side API
  - *clients to utilize WebSocket, supported by Chrome, Safari*
  - *Hides complexity of WebSocket protocol for the developer*
- JavaScript example

Sockets.IO

- Many options for streaming
  - *long-polling, streaming, iframe, WebSockets*
  - *Not all browsers support WebSockets*
  - *Socket.IO – a layer providing a unified API*

- Sockets.IO
  - *API and JavaScript implementation*
  - *checks the availability of WebSocket protocol*
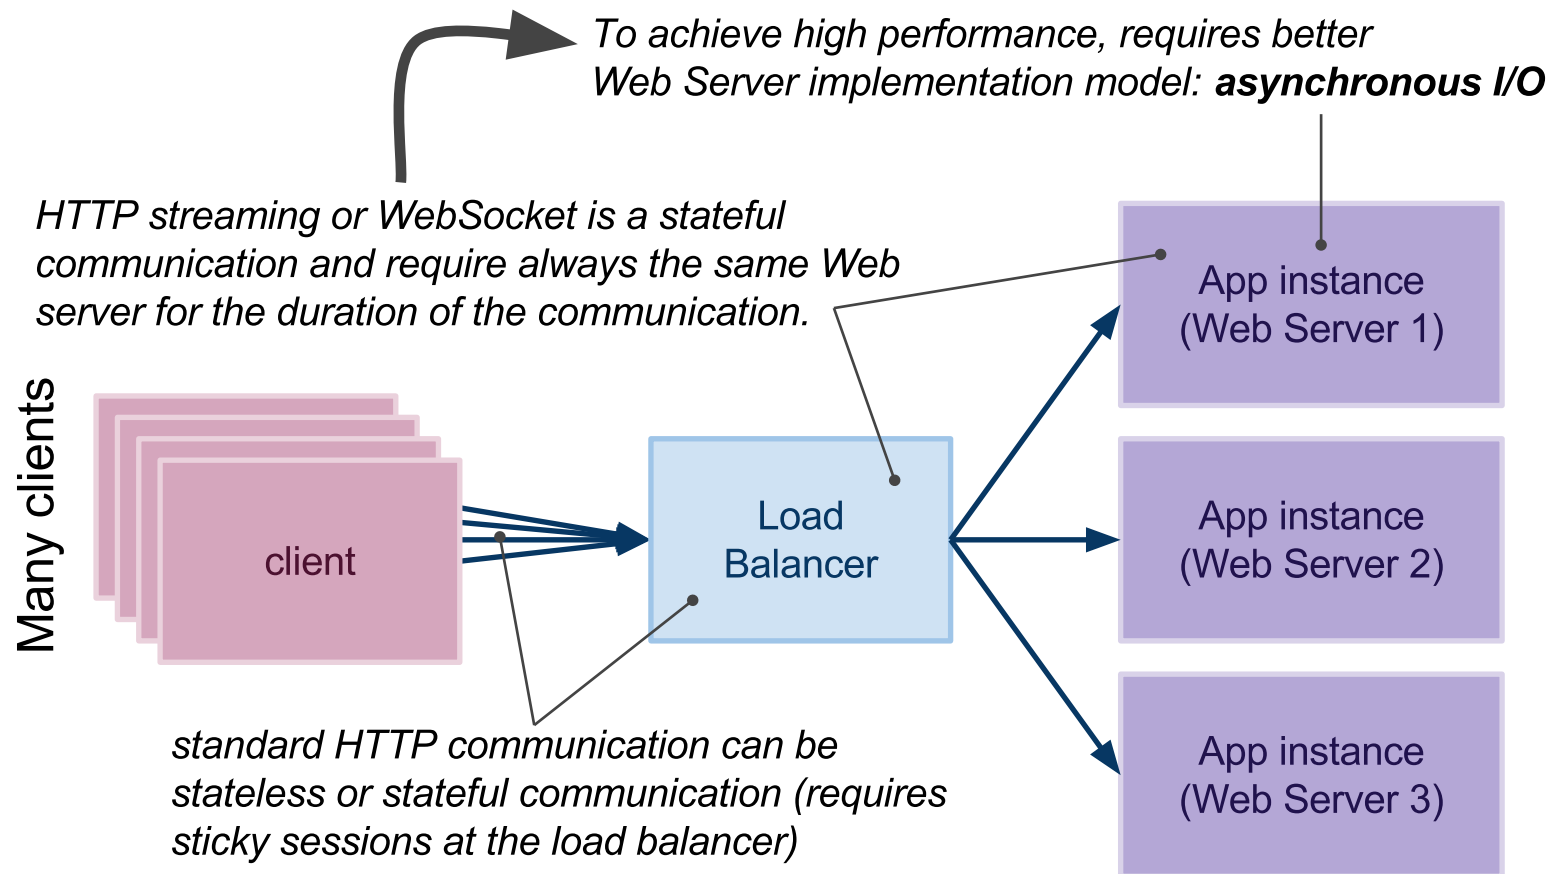   - → *fallback to long-polling or other technologies when not available*

Overview

- Long-polling and Streaming
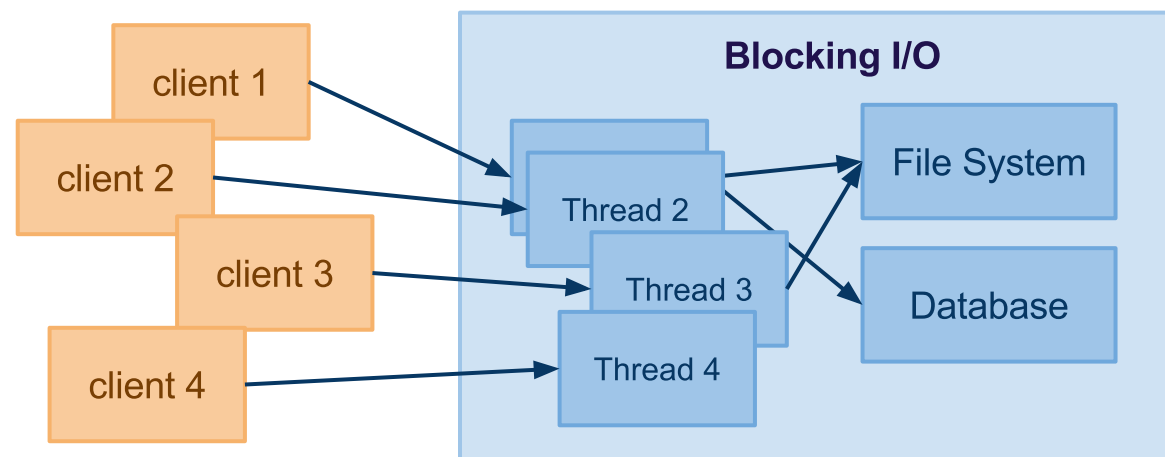- WebSocket Protocol
- New I/O Model

Highly Scalable Web Servers

- Concurrent connections
  - *servers must serve a huge amount of concurrent connections*
  - *Highly scalable Web apps*
    - → *many concurrent requests at the same time*
    - → *QPS: 10-100 or more (GAE scales up to 500 QPS)*
  - *more significant with new trends regarding streaming (HTTP and WebSocket)*

- Web server implementation models:
  Synchronous I/O vs. Asynchronous I/O
  - *synchronous I/O (aka blocking I/O)*
    - → *traditional: server creates a thread for every connection*
  - *asynchronous I/O (aka non-blocking I/O)*
    - → *new one, server handles processing of requests separately from incoming connections*

Web App Scalability

To achieve high performance, requires better
Web Server implementation model: **asynchronous I/O**

*HTTP streaming or WebSocket is a stateful
communication and require always the same Web
server for the duration of the communication.*

Many clients

client

Load
Balancer

App instance
(Web Server 1)

App instance
(Web Server 2)

App instance
(Web Server 3)

*standard HTTP communication can be
stateless or stateful communication (requires
sticky sessions at the load balancer)*

Synchronous I/O Model

- every request served by a single thread
  - *reserved for the whole processing, the thread is "blocked"*
- when processing of the request is fast, scales well
  - *OS maintains a pool of threads that are reused for new requests*
- when processing of the request requires other interactions with DB/FS or network communication is slow → scaling is bad
  - *more significant with streaming (long polling or HTTP streaming)*
- OS may create couple of hundreds of threads (~1000 is very large)
 - *app may serve over 1K clients easily*

Asynchronous I/O Model
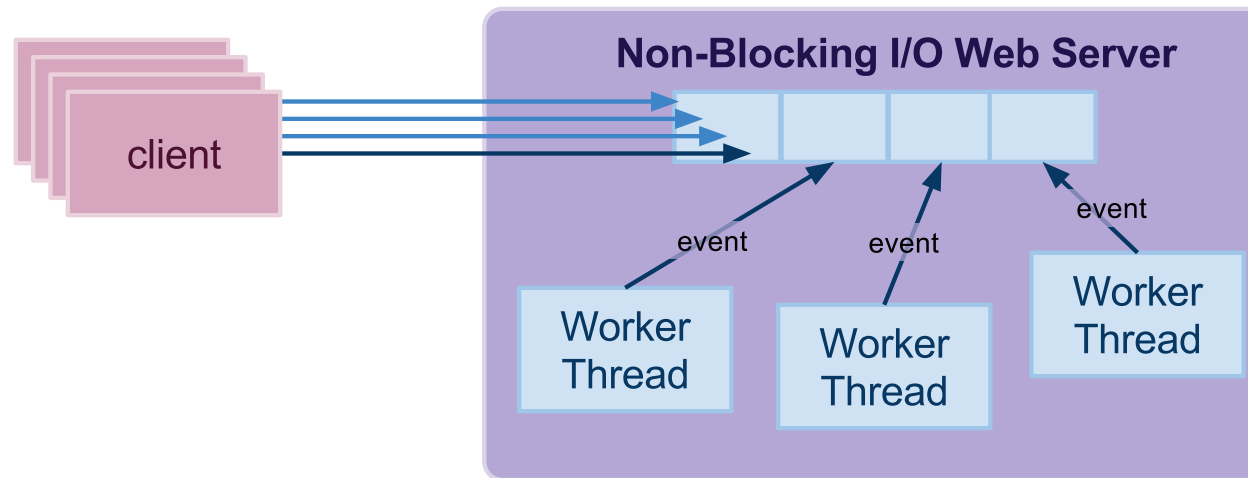
- requests/connections maintained by the OS

- Web server reacts on the events
  - *such as new socket, read, write*
  - *it may create a working thread to perform required processing*
  - *Web server may control the number of Worker Threads*

- significantly less number of working threads as opposed to blocking I/O

Performance Experiment
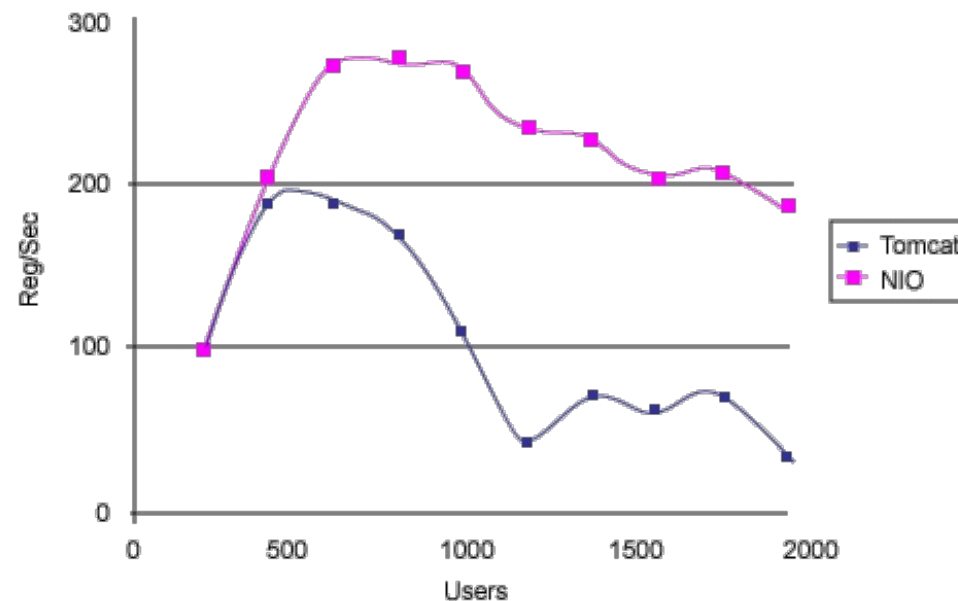
Non-blocking vs. blocking performance (number of requests per second served by the server vs. number of users), source The Servlet API and NIO: Together at last 🔗

- Tomcat – Java-based, uses I/O blocking communication
  - *configured to run up to 2,000 threads*
- NIO – a Web server implemented using Java.NIO (Java New I/O)
  - *only 4 working threads*
- simple HTTP `GET` serving textual content

Emerging Technologies

- Node.js
  - *NodeJS* ⧉ *– event-driven I/O framework on JavaScript V8 engine*
    → every I/O as event:
  - *runs in Linux/Unix/OS X environments*
  - *Executes your server-side JavaScript code*
  - *Socket.IO as a modul provides a streaming layer*

- Java.NIO
  - *Java New I/O, standard in Java SE 7*