

Web 2.0

Lecture 8: Protocols for the Realtime Web

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <http://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <http://vitvar.com/courses/w20>



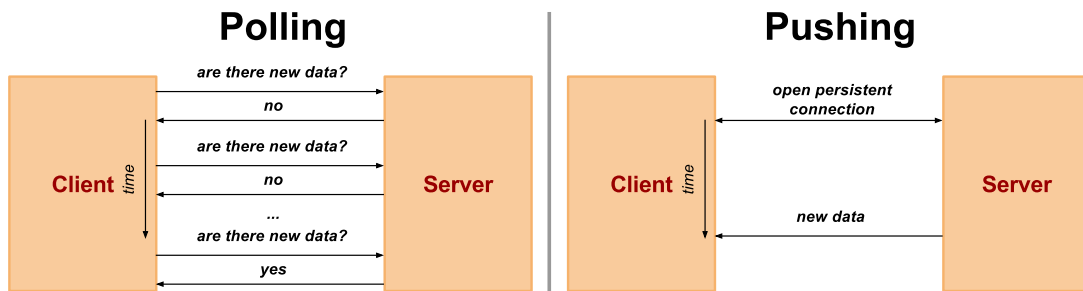
Evropský sociální fond
Praha & EU: Investujeme do vaší budoucnosti

Modified: Mon Apr 14 2014, 09:11:46
Humla v0.3

Overview

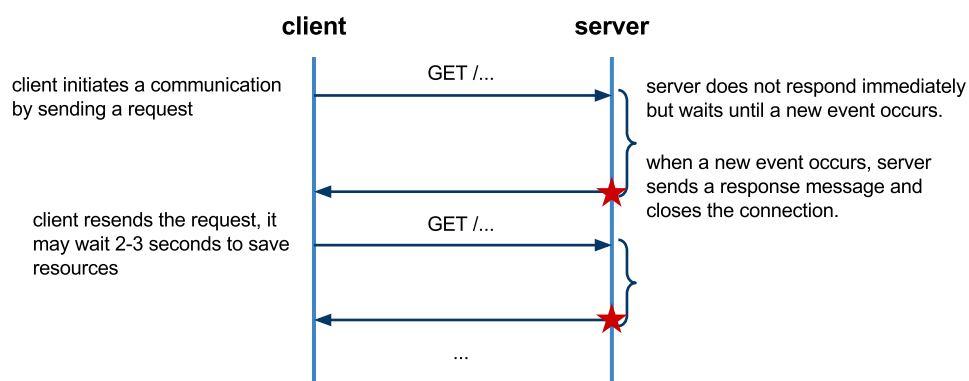
- Long-polling and Streaming
- WebSocket Protocol
- New I/O Model

Pushing and Polling



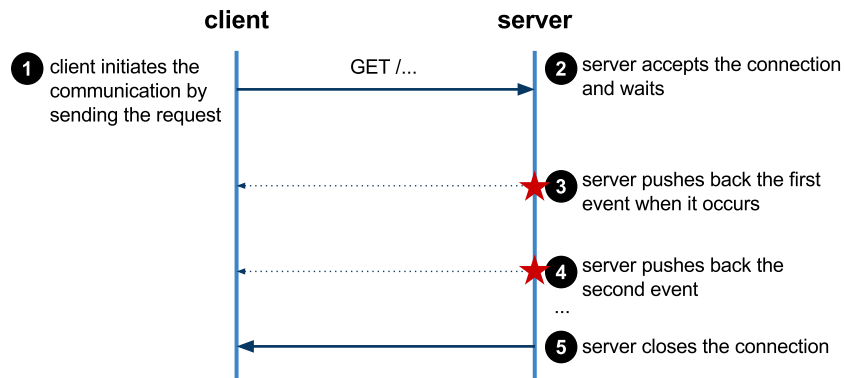
- Conceptual basis in messaging architectures
 - event-driven architectures (EDA)
- **HTTP is a request-response protocol**
 - response cannot be sent without request
 - server cannot initiate the communication
- **Polling** – client periodically checks for updates on the server
- **Pushing** – updates from the server (also called COMET)
 - = **long polling** – server holds the request for some time
 - = **streaming** – server sends updates without closing the socket

HTTP Long Polling



- Server holds long-poll requests
 - server responds when an event or a timeout occurs
 - saves computing resources at the server as well as network resources
 - can be applied over HTTP persistent and non-persistent communication
- Issues:
 - maximum time of the request processing at the server
 - concurrent requests processing at the server

HTTP Streaming



- server defers the response until an event or timeout is available
- when an event is available, server sends it back to client as part of the response; this does not terminate the connection
- server is able to send pieces of response w/o terminating the conn.
 - using **transfer-encoding** header in HTTP 1.1
 - using *End of File* in HTTP 1.0
(server omits **content-length** in the response)

Chunked Response

- Transfer encoding **chunked**
 - It allows to send multiple sets of data over a single connection
 - a chunk represents data for the event
 - Each chunk starts with hexadecimal value for length
 - End of response is marked with the chunk length of 0
- Steps:
 - server sends HTTP headers and the first chunk (step 3)
 - server sends second and subsequent chunk of data (step 4)
 - server terminates the connection (step 5)

Issues with Chunked Response

- Chunks vs. Events
 - *chunks cannot be considered as app messages (events)*
 - *intermediaries might "re-chunk" the message stream*
 - *e.g., combining different chunks into a longer one*
- Client Buffering
 - *clients may buffer all data chunks before they make the response available to the client application*
- HTTP streaming in browsers
 - *Server-sent events*

Server-Sent Events

- W3C specification
 - *part of HTML5 specs, see*
 - *API to handle HTTP streaming in browsers by using DOM events*
 - *transparent to underlying HTTP streaming mechanism*
 - *can use both chunked messages and EOF*
 - *same origin policy applies*
- **EventSource** interface
 - *event handlers: **onopen**, **onmessage**, **onerror***
 - *constructor **EventSource(url)** – creates and opens the stream*
 - *method **close()** – closes the connection*
 - *attribute **readyState***
 - **CONNECTING** – *The connection has not yet been established, or it was closed and the user agent is reconnecting.*
 - **OPEN** – *The user agent has an open connection and is dispatching events as it receives them.*
 - **CLOSED** – *The conn. is not open, the user agent is not reconnecting.*

Example

- Initiating **EventSource**
- Defining event handlers
 - *when the conn. is closed, the browser reconnects every ~3 seconds*
→ can be changed using **retry** attribute in the message data

Event Stream Format

- Format
 - *response's **content-type** must be **text/event-stream***
 - *every line starts with **data:**, event message terminates with 2 \n chars.*
 - *every message may have associated **id** (is optional)*
- JSON data in multiple lines of the message
- Changing the reconnection time
 - *default is 3 seconds*

Server-side implementation

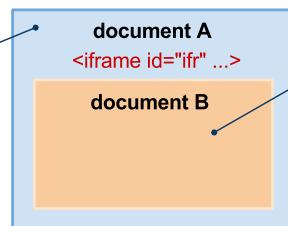
- Java Servlet
 - *method* **doGet**

Other Technologies

- Cross-document messaging

script in document A

```
var o=document.getElementById("ifr");  
o.contentWindow.postMessage("Hello world",  
    "http://example.org/")
```



script in document B

```
window.addEventListener('message', receiver, false);  
function receiver(e) {  
    if (e.origin == 'http://example.com') {  
        if (e.data == 'Hello world') {  
            e.source.postMessage('Hello', e.origin);  
        } else {  
            alert(e.data);  
        }  
    }  
}
```

- *The use of Cross Document Messaging for streaming*
 1. The client loads a streaming resource in a hidden **iframe**
 2. The server pushes a JavaScript code to the **iframe**
 3. The browser executes the code as it arrives from the server
 4. The embedded iframe's code posts a message to the upper document

- Channel API

- Google Technology for streaming API for AppEngine
- not based on HTTP streaming
- utilizes XMPP capabilities + hidden iframe at client-side

Overview

- Long-polling and Streaming
- **WebSocket Protocol**
- New I/O Model

WebSocket

- Specifications
 - *IETF defines*
 - *W3C defines*
- Design principles
 - *a new protocol*
 - *browsers, web servers, and proxy servers need to support it*
 - *a layer on top of TCP*
 - *bi-directional communication between client and servers*
 - *low-latency apps without HTTP overhead*
 - *Web origin-based security model for browsers*
 - *same origin policy, cross-origin resource sharing*
 - *support multiple server-side endpoints*
- Two phases
 - *Handshake – as an **upgrade** of a HTTP connection*
 - *data transfer – the protocol-specific on-the-wire data transfer*

Handshake – Request

- Request
 - *client sends a following HTTP request to upgrade the connection to WebSocket*
 - **Connection** – *request to upgrade the protocol*
 - **Upgrade** – *protocol to upgrade to*
 - **Sec-WebSocket-Key** – *a client key for later validation*
 - **Sec-WebSocket-Origin** – *origin of the request*
 - **Sec-WebSocket-Protocol** – *list of sub-protocols that client supports (proprietary)*

Handshake – Response

- Response
 - *server accepts the request and responds as follows*
 - **101 Switching Protocols** – *status code for a successful upgrade*
 - **Sec-WebSocket-Protocol** – *a sub-protocol that the server selected from the list of protocols in the request*
 - **Sec-WebSocket-Accept** – *a key to prove it has received a client WebSocket handshake request*
 - *Formula to compute Sec-WebSocket-Accept*
 - **SHA-1** – *hashing function*
 - **Base64Encode** – *Base64 encoding function*
 - **"258EAF55-E914-47DA-95CA-C5AB0DC85B11"** – *magic number*

Data Transfer

- After successful handshake
 - *socket between the client and the "resource" at the server is established*
 - *client and the server can both read and write from/to the socket*
 - *No HTTP headers overhead*
- Data Framing
 - *defines a format for data transmitted in TCP packets*
 - *payload length, closing frame, ping, pong, type of data (text/binary), etc. and payload (message data)*

WebSocket API

- Client-side API
 - *clients to utilize WebSocket, supported by Chrome, Safari*
 - *Hides complexity of WebSocket protocol for the developer*
- JavaScript example

Sockets.IO

- Many options for streaming
 - *long-polling, streaming, iframe, WebSockets*
 - *Not all browsers support WebSockets*
 - *– a layer providing a unified API*
- Sockets.IO
 - *API and JavaScript implementation*
 - *checks the availability of WebSocket protocol*
 - *fallback to long-polling or other technologies when not available*

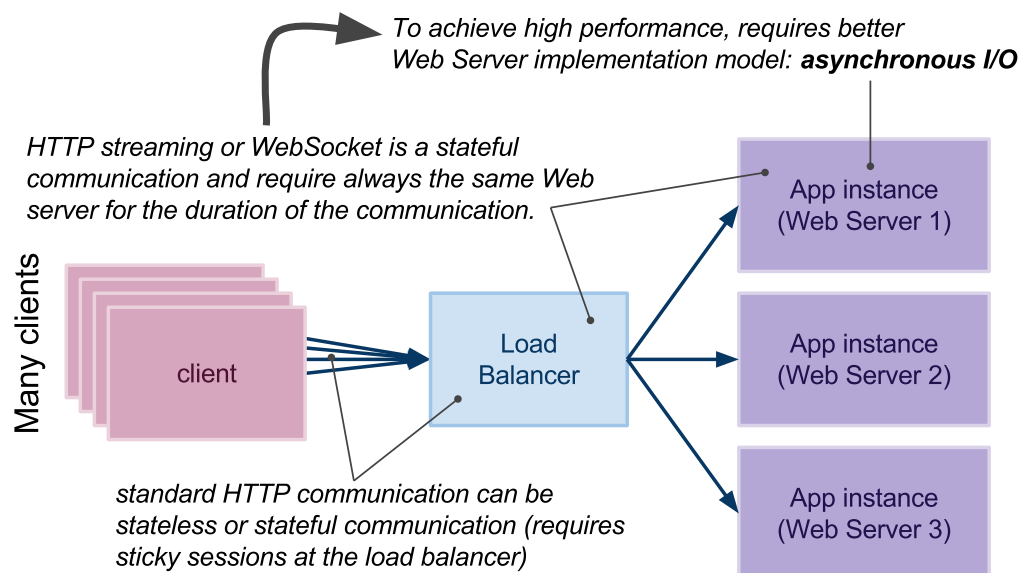
Overview

- Long-polling and Streaming
- WebSocket Protocol
- **New I/O Model**

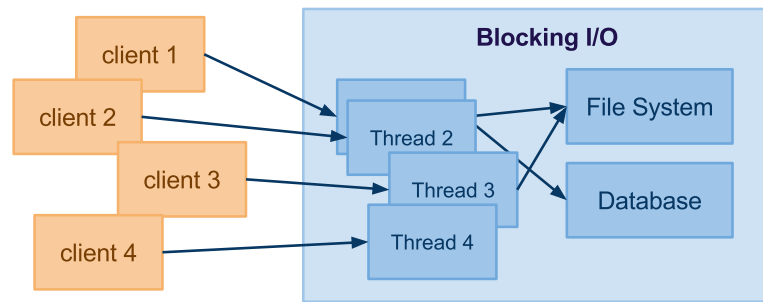
Highly Scalable Web Servers

- Concurrent connections
 - servers must serve a huge amount of concurrent connections
 - Highly scalable Web apps
 - many concurrent requests at the same time
 - QPS: 10-100 or more (GAE scales up to 500 QPS)
 - more significant with new trends regarding streaming (HTTP and WebSocket)
- Web server implementation models:
 - Synchronous I/O vs. Asynchronous I/O**
 - synchronous I/O (aka blocking I/O)
 - traditional: server creates a thread for every connection
 - asynchronous I/O (aka non-blocking I/O)
 - new one, server handles processing of requests separately from incoming connections

Web App Scalability

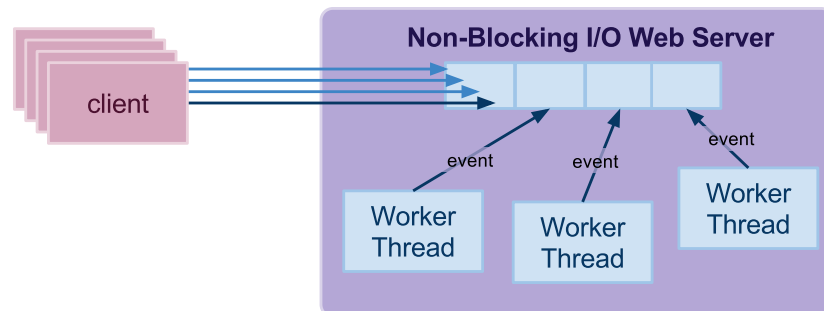


Synchronous I/O Model



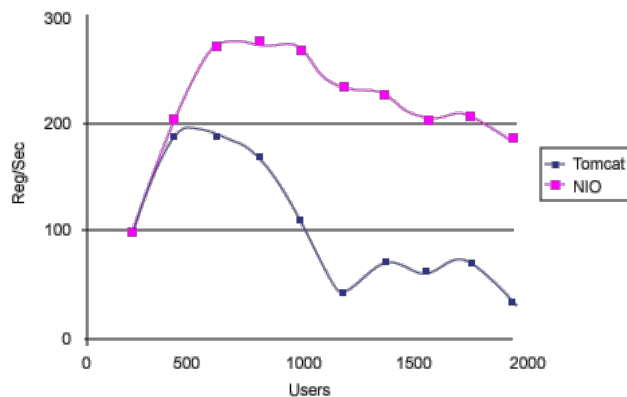
- every request served by a single thread
 - reserved for the whole processing, the thread is "blocked"
- when processing of the request is fast, scales well
 - OS maintains a pool of threads that are reused for new requests
- when processing of the request requires other interactions with DB/FS or network communication is slow → scaling is bad
 - more significant with streaming (long polling or HTTP streaming)
- OS may create couple of hundreds of threads (~1000 is very large)
 - can't serve over 1K clients easily

Asynchronous I/O Model



- requests/connections maintained by the OS
- Web server reacts on the events
 - such as new socket, read, write
 - it may create a working thread to perform required processing
 - Web server may control the number of Worker Threads
- significantly less number of working threads as opposed to blocking I/O

Performance Experiment



Non-blocking vs. blocking performance (number of requests per second served by the server vs. number of users), source

- Tomcat – Java-based, uses I/O blocking communication
 - *configured to run up to 2,000 threads*
- NIO – a Web server implemented using Java.NIO (Java New I/O)
 - *only 4 working threads*
- simple HTTP **GET** serving textual content

Emerging Technologies

- Node.js
 - *event-driven I/O framework on JavaScript V8 engine*
 - *every I/O as event:*
 - *runs in Linux/Unix/OS X environments*
 - *Executes your server-side JavaScript code*
 - *Socket.IO as a modul provides a streaming layer*
- Java.NIO
 - *Java New I/O, standard in Java SE 7*